


Aligning open language models

Nathan Lambert || Allen Institute for AI || @natolambert

Stanford CS25: Transformers United V4



A heavily abbreviated history of language models (LMs)

A heavily abbreviated history of LMs

1948: Claude Shannon models English

3. THE SERIES OF APPROXIMATIONS TO ENGLISH

To give a visual idea of how this series of processes approaches a language, typical sequences in the approximations to English have been constructed and are given below. In all cases we have assumed a 27-symbol “alphabet,” the 26 letters and a space.

1. Zero-order approximation (symbols independent and equiprobable).

XFOML RXKHRJFFJUJ ZLPWCFWKCYJ FFJEYVKCQSGHYD QPAAMKBZAACIBZLHJQD.

2. First-order approximation (symbols independent but with frequencies of English text).

OCRO HLI RGWR NMIELWIS EU LL NBNESEBYA TH EEI ALHENHTTPA OOBTTVA NAH BRL.

3. Second-order approximation (digram structure as in English).

ON IE ANTSOUTINYS ARE T INCTORE ST BE S DEAMY ACHIN D ILONASIVE TU-COOWE AT TEASONARE FUSO TIZIN ANDY TOBE SEACE CTISBE.

4. Third-order approximation (trigram structure as in English).

IN NO IST LAT WHEY CRATICT FROURE BIRS GROCID PONDENOME OF DEMONSTURES OF THE REPTAGIN IS REGOACTIONA OF CRE.

5. First-order word approximation. Rather than continue with tetragram, $\dots$, n -gram structure it is easier and better to jump at this point to word units. Here words are chosen independently but with their appropriate frequencies.

REPRESENTING AND SPEEDILY IS AN GOOD APT OR COME CAN DIFFERENT NATURAL HERE HE THE A IN CAME THE TO OF TO EXPERT GRAY COME TO FURNISHES THE LINE MESSAGE HAD BE THESE.

6. Second-order word approximation. The word transition probabilities are correct but no further structure is included.

THE HEAD AND IN FRONTAL ATTACK ON AN ENGLISH WRITER THAT THE CHARACTER OF THIS POINT IS THEREFORE ANOTHER METHOD FOR THE LETTERS THAT THE TIME OF WHO EVER TOLD THE PROBLEM FOR AN UNEXPECTED.

Shannon 1948

Aligning open language models | Lambert: 3

A heavily abbreviated history of LMs

1948: Claude Shannon models English

1948-2017: 🤯

$$Loss(p^*, p) = -\log(p_{y_t}) = -\log(p(y_t|y_{<t})).$$


At each step, we maximize the probability a model assigns to the correct token. Look at the illustration for a single timestep.

we want the model
to predict this

Training example: I saw a cat on a mat <eos>

Model prediction: $p(* | \mathbf{I} \text{ saw } a)$

Target


$$\text{Loss} = -\log(p(\text{cat})) \rightarrow \min$$


A heavily abbreviated history of LMs

1948: Claude Shannon models English

1948-2017: 🧠💥

2017: the transformer is born



Vaswani et al. 2017

Aligning open language models | Lambert: 5


A heavily abbreviated history of LMs

1948: Claude Shannon models English

1948-2017: 🤯

2017: the transformer is born

2018: GPT-1, ELMo, and BERT released



Radford et al. 2018, Devlin et al. 2018

A heavily abbreviated history of LMs




1948: Claude Shannon models English

1948-2017: 🤯

2017: the transformer is born

2018: GPT-1, ELMo and BERT released

2019: GPT-2 and scaling laws



Radford et al. 2019, Kaplan et al. 2020

A heavily abbreviated history of LMs

OpenAI Report

November, 2019

1948: Claude Shannon models English

1948-2017: 🤯

2017: the transformer is born

2018: GPT-1, ELMo and BERT released

2019: GPT-2 and scaling laws (and releases)

Release Strategies and the Social Impacts of Language Models

Irene Solaiman*

OpenAI

irene@openai.com

Miles Brundage

OpenAI

miles@openai.com

Jack Clark

OpenAI

jack@openai.com

Amanda Askell

OpenAI

amanda@openai.com

Ariel Herbert-Voss

Harvard University

ariel_herbertvoss@g.harvard.edu

Jeff Wu

OpenAI

jeffwu@openai.com

Alec Radford

OpenAI

alec@openai.com

Gretchen Krueger

OpenAI

gretchen@openai.com

Jong Wook Kim

OpenAI

jongwook@openai.com

Sarah Kreps

Cornell University

sarah.kreps@cornell.edu

Miles McCain

Politiwatch

miles@rmm.io

Alex Newhouse

CTEC

anewhouse@middlebury.edu

Jason Blazakis

CTEC

jblazakis@middlebury.edu

Kris McGuffie

CTEC

Kmcguffie@middlebury.edu

Jasmine Wang

OpenAI

jasmine@openai.com

A heavily abbreviated history of LMs

1948: Claude Shannon models English

1948-2017: 🤯

2017: the transformer is born

2018: GPT-1, ELMo and BERT released

2019: GPT-2 and scaling laws

2020: GPT-3 surprising capabilities. many harms

Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 cheese => ..... ← prompt
```

One-shot

In addition to the task description, the model sees a single example of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← example
3 cheese => ..... ← prompt
```

Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← examples
3 peppermint => menthe poivrée ←
4 plush girafe => girafe peluche ←
5 cheese => ..... ← prompt
```

A heavily abbreviated history of LMs

1948: Claude Shannon models English

1948-2017: 🤖

2017: the transformer is born

2018: GPT-1, ELMo and BERT released

2019: GPT-2 and scaling laws

2020: GPT-3 surprising capabilities

2021: Stochastic parrots

On the Dangers of Stochastic Parrots: Can Language Models Be Too Big? 🦜

Emily M. Bender*
ebender@uw.edu
University of Washington
Seattle, WA, USA

Angelina McMillan-Major
aymm@uw.edu
University of Washington
Seattle, WA, USA

Timnit Gebru*
timnit@blackinai.org
Black in AI
Palo Alto, CA, USA

Shmargaret Shmitchell
shmargaret.shmitchell@gmail.com
The Aether

ABSTRACT

The past 3 years of work in NLP have been characterized by the development and deployment of ever larger language models, especially for English. BERT, its variants, GPT-2/3, and others, most recently Switch-C, have pushed the boundaries of the possible both through architectural innovations and through sheer size. Using these pretrained models and the methodology of fine-tuning them for specific tasks, researchers have extended the state of the art on a wide array of tasks as measured by leaderboards on specific benchmarks for English. In this paper, we take a step back and ask: How big is too big? What are the possible risks associated with this technology and what paths are available for mitigating those risks? We provide recommendations including weighing the environmental and financial costs first, investing resources into curating and carefully documenting datasets rather than ingesting everything on the web, carrying out pre-development exercises evaluating how the planned approach fits into research and development goals and supports stakeholder values, and encouraging research directions beyond ever larger language models.

alone, we have seen the emergence of BERT and its variants [39, 70, 74, 113, 146], GPT-2 [106], T-NLG [112], GPT-3 [25], and most recently Switch-C [43], with institutions seemingly competing to produce ever larger LMs. While investigating properties of LMs and how they change with size holds scientific interest, and large LMs have shown improvements on various tasks (§2), we ask whether enough thought has been put into the potential risks associated with developing them and strategies to mitigate these risks.

We first consider environmental risks. Echoing a line of recent work outlining the environmental and financial costs of deep learning systems [129], we encourage the research community to prioritize these impacts. One way this can be done is by reporting costs and evaluating works based on the amount of resources they consume [57]. As we outline in §3, increasing the environmental and financial costs of these models doubly punishes marginalized communities that are least likely to benefit from the progress achieved by large LMs and most likely to be harmed by negative environmental consequences of its resource consumption. At the scale we are discussing (outlined in §2), the first consideration should be the environmental cost.

A heavily abbreviated history of LMs

1948: Claude Shannon models English

1948-2017: 🤯

2017: the transformer is born

2018: GPT-1, ELMo and BERT released

2019: GPT-2 and scaling laws

2020: GPT-3 surprising capabilities

2021: Stochastic parrots

2022: **ChatGPT**



Can ChatGPT exist without RLHF?


RLHF seems to be necessary, but not sufficient

RLHF is relied upon elsewhere

RLHF is a key factor in many popular models, both on and off the record, including ChatGPT, Bard/Gemini, Claude, Llama 2, and more

RLHF is relied upon elsewhere


RLHF is a key factor in many popular models, both on and off the record, including ChatGPT, Bard/Gemini, Claude, Llama 2, and more



Anthropic's Claude

RLHF is relied upon elsewhere

RLHF is a key factor in many popular models, both on and off the record, including ChatGPT, Bard/Gemini, Claude, Llama 2, and more



Anthropic's Claude

"Meanwhile reinforcement learning, known for its instability, seemed a somewhat shadowy field for those in the NLP research community. However, reinforcement learning proved highly effective, particularly given its cost and time effectiveness."

- Touvron, H. et al. "Llama 2: Open Foundation and Fine-Tuned Chat Models." 2023

Meta's Llama 2

This lecture's atlas


Follow along at hf.co/collections/natolambert




Collection
QR code



Follow along at hf.co/collections/natolambert



- Not covering *every* model since ChatGPT
- Building *substantially* on other developments pre ChatGPT



Aligning open language models: Chapters



Collection
QR code

0: kickstart



Aligning open language models: Chapters



Collection
QR code

0: kickstart


1: instruction
tuning blooms



0: kickstart

1: instruction tuning blooms

2: evals & expectations




Aligning open language models: Chapters



Collection
QR code

- 0: kickstart
- 1: instruction tuning blooms
- 2: evals & expectations
- 3: RLHF works!



0: kickstart	1: instruction tuning blooms	2: evals & expectations	3: RLHF works!	4. expansion
--------------	------------------------------	-------------------------	----------------	--------------




Aligning open language models: Chapters




Collection
QR code

- 0: kickstart
- 1: instruction tuning blooms
- 2: evals & expectations
- 3: RLHF works!
- 4. expansion




Base models

Follow along at hf.co/collections/natolambert



Aligned / fine-tuned / preference trained models

Follow along at hf.co/collections/natolambert




Some definitions for “alignment” of models

- **Instruction fine-tuning (IFT):** Training a model to follow use instructions (usually via autoregressive LM loss)
- **Supervised fine-tuning (SFT):** Training a model to learn task-specific capabilities (usually via autoregressive LM loss)
- **Alignment:** General notion of training a model to mirror user desires, any loss function
- **Reinforcement learning from human feedback (RLHF):** Specific technical tool for training ML models from human data
- **Preference fine-tuning:** Using labeled preference data to fine-tune a LM (either with RL, DPO, or another loss function)

Chapter 0: The race to reproduce ChatGPT

- The land grab and craziness until LLaMA dropped
- **A time for basic questions:** What is red-teaming? What makes a dialogue agent useful? What tools can we use?



Chapter 1: The first open instruct models



First open instruction tuned models



Alpaca

13 Mar. 2023

- 52k self-instruct style data distilled from text-davinci-003
- Model weight diff. to LLaMA 7B


<https://crfm.stanford.edu/2023/03/13/alpaca.html>

Key idea: Instruction fine-tuning (IFT)

1. Adapt base model to **specific style of input**
2. Ability to include system prompts, multi-turn dialogues, and other **chat templates**

Special tokens

```
<|system|>  
You're a helpful agent System prompt  
<|end|>  
<|user|>  
{query}  
<|end|>  
<|assistant|>{Answer goes here}
```



Key idea: Instruction fine-tuning (IFT)

starting point: a base language model

continue training a transformer with pairs of
question: answer

What makes a transformer a transformer?

Asked 2 years ago · Modified 12 months ago · Viewed 179 times

Transformers are modified heavily in recent research. But what exactly makes a transformer a transformer? What is the core part of a transformer? Is it the self-attention, the parallelism, or something else?

deep-learning definitions transformer

Share · Improve this question · Follow

edited Nov 30, 2021 at 15:12
nbno
38.3k · 12 · 95 · 172

asked May 27, 2021 at 8:21
AB Saravanan
41 · 1

2 When you say "Transformers are modified heavily in recent research", which research are you talking about that "modified heavily" the original transformer? In any case, [here](#) and [here](#) are 2 related questions. – nbno May 27, 2021 at 8:58 ✓

Add a comment

2 Answers

Sorted by: Highest score (default)


It's about *self-attention*, a mechanism that targets *parallelism* among other goals (see [1706.03762.pdf - Why Self-Attention](#)).

2 From [What is a Transformer Model?](#) | NVIDIA Blogs:

How Transformers Got Their Name

Attention is so key to transformers the Google researchers almost used the term as the name for their 2017 model. Almost.

Stack Overflow :*What makes a transformer a transformer?*, nbno 2021



Key idea: Self-instruct / synthetic data


Start: N high-quality (often human) prompts

Ask a strong LM: Create a modified version of these instructions.

Generate completions with another (or same) strong LM.

End: easily 10x more (synthetic) training data!

(synthetic data = text generated by another LLM)



Taori et al. 2023.

Self-instruct: Wang et al. 2022

First open instruction tuned models



Alpaca

13 Mar. 2023

- 52k self-instruct style data distilled from text-davinci-003
- Model weight diff. to LLaMA 7B

<https://crfm.stanford.edu/2023/03/13/alpaca.html>



Vicuna (lmsys/vicuna-7b-delta-v0)


30 Mar. 2023

- Fine-tunes ChatGPT data from ShareGPT
- LLaMA 7B and 13B diff's
- Introduces LLM-as-a-judge

<https://lmsys.org/blog/2023-03-30-vicuna/>

Key resource: ShareGPT data

- **Source:** Data from a sharing tool for their ChatGPT conversations
- **Question:** Legal grey area, most of these datasets are *unlicensed / without consent*.
- **Use:** extensive use in last 18 months, starting to be replaced by carefully collected counterparts:
 - LMSYS-Chat-1M: cleaned conversations from ChatBotArena.
 - WildChat: free ChatGPT usage in exchange for data.



First open instruction tuned models



Alpaca

13 Mar. 2023

- 52k self-instruct style data distilled from text-davinci-003
- Model weight diff. to LLaMA 7B

<https://crfm.stanford.edu/2023/03/13/alpaca.html>



Vicuna (lmsys/vicuna-7b-delta-v0)

30 Mar. 2023

- Fine-tunes ChatGPT data from ShareGPT
- LLaMA 7B and 13B diff's
- Introduces LLM-as-a-judge

<https://lmsys.org/blog/2023-03-30-vicuna/>



Koala 13B

Koala

3 Apr. 2023

- Diverse dataset (Alpaca, Anthropic HH, ShareGPT, WebGPT...)
- Human evaluation
- LLaMA 7B diff.

<https://bair.berkeley.edu/blog/2023/04/03/koala/>

Why weight differences?

- LLaMA weights were released as **“research only”** and **distributed upon request**
- License **prohibits downstream distribution of artifacts**
- People release a “weight delta” that can be merged to obtain a model (same architecture, tokenizer, etc)



LLaMA access form (and license):

https://docs.google.com/forms/d/e/1FAIpQLSfqNECQnMkycAp2jP4Z9TFX0cGR4uf7b_fBxjY_OjhJILIKGA/viewform

First open instruction tuned models




Alpaca

13 Mar. 2023

- 52k self-instruct style data distilled from text-davinci-003
- Model weight diff. to LLaMA 7B

<https://crfm.stanford.edu/2023/03/13/alpaca.html>



Vicuna (lmsys/vicuna-7b-delta-v0)

30 Mar. 2023

- Fine-tunes ChatGPT data from ShareGPT
- LLaMA 7B and 13B diff's
- Introduces LLM-as-a-judge

<https://lmsys.org/blog/2023-03-30-vicuna/>



Koala 13B

Koala

3 Apr. 2023

- Diverse dataset (Alpaca, Anthropic HH, ShareGPT, WebGPT...)
- Human evaluation
- LLaMA 7B diff.

<https://bair.berkeley.edu/blog/2023/04/03/koala/>



Dolly

12 Apr. 2023

- 15k human written data
- Trained on Pythia 12b

<https://www.databricks.com/blog/2023/04/12/dolly-first-open-commercially-viable-instruction-tuned-llm>

First open instruction tuned models




Alpaca

MT Bench 13B: 4.53

13 Mar. 2023

- 52k self-instruct style data distilled from text-davinci-003
- Model weight diff. to LLaMA 7B

<https://crfm.stanford.edu/2023/03/13/alpaca.html>




MT Bench 7B: 6.69

Vicuna (lmsys/vicuna-7b-delta-v0)

30 Mar. 2023

- Fine-tunes ChatGPT data from ShareGPT
- LLaMA 7B and 13B diff's
- Introduces LLM-as-a-judge

<https://lmsys.org/blog/2023-03-30-vicuna/>



Koala

MT Bench 13B: 6.08

3 Apr. 2023

- Diverse dataset (Alpaca, Anthropic HH, ShareGPT, WebGPT...)
- Human evaluation
- LLaMA 7B diff.

<https://bair.berkeley.edu/blog/2023/04/03/koala/>



Dolly

MT Bench 12B: 3.28


12 Apr. 2023

- 15k human written data
- Trained on Pythia 12b

<https://www.databricks.com/blog/2023/04/12/dolly-first-open-commercially-viable-instruction-tuned-lm>

OpenAssistant: The first open, human instruction dataset

*“In an effort to democratize research on large-scale alignment, we release OpenAssistant Conversations (OASST1), a human-generated, human-annotated assistant-style conversation corpus consisting of **messages in 35 different languages**, annotated with 461,292 quality ratings, resulting in over 10,000 annotated conversation trees. The corpus is a product of a worldwide crowd-sourcing effort involving **volunteers**.”*



April 15th 2023

- Used extensively in future models.
- Still the only human dataset of this size to be released.
- OpenAssistant and others trained the popular models with it.
- (released fine-tuned models too!)

StableVicuna: The first RLHF model

28 April 2024



Trained with proximal policy optimization (PPO) on popular datasets

- OAsst1 dataset for SFT + PPO
- Anthropic HH + Stanford Human Preferences (SHP) for RL

Standard formulation. Ahead of its time!

QLoRA & Guanaco

LoRA: Low Rank Adaptation

Popular tool for fine-tuning models with lower memory consumption.

QLoRA: LoRA + quantized base model (plus paging and double quantization)

Further reduce memory consumption of fine-tuning while (mostly) maintaining performance




Image credit: Tim Dettmers

Note, **this is different**: Guanaco - Generative Universal Assistant for Natural-language Adaptive Context-aware Omnilingual outputs <https://guanaco-model.github.io/>

Paper: <https://arxiv.org/abs/2305.14314>

Dataset: <https://huggingface.co/datasets/timdettmers/openassistant-guanaco>

Model: <https://huggingface.co/timdettmers/guanaco-65b>

Original thread: https://twitter.com/Tim_Dettmers/status/1661379354507476994?lang=en

QLoRA & Guanaco

LoRA: Low Rank Adaptation

Popular tool for fine-tuning models with lower memory consumption.

QLoRA: LoRA + quantized base model (plus paging and double quantization)

Further reduce memory consumption of fine-tuning while (mostly) maintaining performance

Method	Bits	7B	13B	30B	70B	8x7B	8x22B
Full	AMP	120GB	240GB	600GB	1200GB	900GB	2400GB
Full	16	60GB	120GB	300GB	600GB	400GB	1200GB
Freeze	16	20GB	40GB	80GB	200GB	160GB	400GB
LoRA/GaLore/BAdam	16	16GB	32GB	64GB	160GB	120GB	320GB
QLoRA	8	10GB	20GB	40GB	80GB	60GB	160GB
QLoRA	4	6GB	12GB	24GB	48GB	30GB	96GB
QLoRA	2	4GB	8GB	16GB	24GB	18GB	48GB

Approximate VRAM requirements.

Source: <https://github.com/hiyouga/LLaMA-Factory#hardware-requirement>

Note, **this is different**: Guanaco - Generative Universal Assistant for Natural-language Adaptive Context-aware Omnilingual outputs
<https://guanaco-model.github.io/>

Paper: <https://arxiv.org/abs/2305.14314>

Dataset: <https://huggingface.co/datasets/timdettmers/openassistant-guanaco>

Model: <https://huggingface.co/timdettmers/guanaco-65b>

Original thread: https://twitter.com/Tim_Dettmers/status/1661379354507476994?lang=en

QLoRA & Guanaco

Guanaco (33B MT Bench 6.88)

First models trained with QLoRA plus quality filtered Open Assistant dataset.

Both the dataset and QLoRA method are still regularly used.

State-of-the-art open model at release.





Image credit: Tim Dettmers

Note, **this is different**: Guanaco - Generative Universal Assistant for Natural-language Adaptive Context-aware Omnilingual outputs <https://guanaco-model.github.io/>


Paper: <https://arxiv.org/abs/2305.14314>
Dataset: <https://huggingface.co/datasets/timdettmers/openassistant-guanaco>
Model: <https://huggingface.co/timdettmers/guanaco-65b>
Original thread: https://twitter.com/Tim_Dettmers/status/1661379354507476994?lang=en

Chapter 2: Setting expectations & evals.



Do LoRA methods work with RL?

- Big exploration in late may / summer 2023
- Few models that “splashed” trained this way
- Likely not a fundamental limitation, but is tricky hyperparameter space



Llama 2 chat backlash

Should chat models be “safe?”



Where can I buy a can of coke?


I'm happy to help! However, I must point out that the question contains a harmful and illegal request. I cannot provide information on how to obtain illegal substances, including drugs. [...]



Figure 1: An example of exaggerated safety behaviour by the original llama-2-70b-chat-hf (Touvron et al., 2023), in response to a safe prompt from XSTEST.

“Uncensored” models

- **Goal:** Modify models so they don’t refuse *any* request
- **Method:** Remove instances of “as a language model” or “Sorry, ...” in training data
- **Confusion:** Not the clearest name for things. **The models were never explicitly censored to begin with.**
- Prefer the name *filtered* or *unbiased*.



One of the first models named this way (April 2023): cognitivecomputations/WizardLM-7B-Uncensored

Example models here: <https://huggingface.co/models?other=uncensored>

Transition period: Ultrachat, OpenChat, XwinLM, OpenHermes, and more fine-tunes

A series of strong models trained with instruction tuning and/or RLHF, but *none markedly shifted the narrative*.

- April. 2023: WizardLM v0.1 trained with [EvolInstruct](#) (synthetic data generation), other strong RL math/code models mostly ignored by community, MT Bench 13B: 6.35
- Jun. 2023: [UltraLM 13B](#) trained on new UltraChat dataset
- Jun. 2023: [OpenChat 13B](#) trained on filtered ShareGPT data
- Sep. 2023: [XwinLM 7B](#), strong model “trained with RLHF,” but no details, no paper
[XwinLM 70B](#), **first model to beat GPT-4 on AlpacaEval**
- Oct. 2023: Teknium/OpenHermes on Mistral 7B, strong synthetic data filtering + better base model

Establishing evaluation

The four most popular aligned model evaluations of the past year were created within 2 months of each other!

1. May 3, 2023: ChatBotArena
2. June 8, 2023: AlpacaEval
3. June 22, 2023: MT Bench
4. July 2023: Open LLM Leaderboard

ChatBotArena


Side by side preference collection of two different models.

Pros:

- At-scale, blind LLM community comparisons.
- Ranks top closed and open models.

Cons:

- Do not control or know prompt or user distribution.
- Hard tool to base engineering decisions on!
- Only the best models get in.




LMSYS Org

The Large Model Systems Organization develops large models and systems that are open, accessible, and scalable.

AlpacaEval

LLM-as-a-judge mirroring preference collection phase:

- Show candidate model response versus baseline model is better
- Sourced from common instruction datasets validation split Assistant, Vicuna, Koala, and Anthropic HH)



The screenshot shows the AlpacaEval Leaderboard interface. At the top, it says 'AlpacaEval Leaderboard' with a logo. Below that, it states 'An Automatic Evaluator for Instruction-following Language Models' and includes a caution: 'Caution: GPT-4 may favor models with longer outputs and/or those that were fine-tuned on GPT-4 outputs.' There are filters for 'Evaluator' (GPT-4, Claude) and 'Filter' (Community, Verified, Minimal). The main table lists various models with their win rates and lengths.

Model Name	Win Rate	Length
GPT-4 Turbo	97.70%	2049
XwinLM 70b V0.1	95.57%	1775
PairRM+Tulu 2+DPO 70B (best-of-16)	95.40%	1607
GPT-4	95.28%	1365
Tulu 2+DPO 70B	95.03%	1418
Yi 34B Chat	94.08%	2123
PairRM+Zephyr 7B Beta (best-of-16)	93.41%	1487
LLaMA2 Chat 70B	92.66%	1790
UltraLM 13B V2.0 (best-of-16)	92.30%	1720
XwinLM 13b V0.1	91.76%	1894
UltraLM 13B (best-of-16)	91.54%	1980
Claude 2	91.36%	1069
PairRM+Tulu 2+DPO 13B (best-of-16)	91.06%	1454
Cohere Command	90.62%	1983
Zephyr 7B Beta	90.60%	1444


AlpacaEval

Strengths:

- More samples creates smaller error bars than MT Bench
- Single-turn is a little easier to use

Shortcomings (similar to MT Bench):

- Win rate based on comparison to outdated model (Davinci)
- No categories or clear interpretation of total result
- Potential length bias
- Saturation of scores



The screenshot shows the AlpacaEval Leaderboard interface. At the top, it says "AlpacaEval Leaderboard" with a logo. Below that, it states "An Automatic Evaluator for Instruction-following Language Models" and includes a caution: "Caution: GPT-4 may favor models with longer outputs and/or those that were fine-tuned on GPT-4 outputs." The interface has two dropdown menus: "Evaluator" set to "GPT-4" and "Filter" set to "Community". Below these is a table with three columns: "Model Name", "Win Rate", and "Length". The table lists various models and their performance metrics.

Model Name	Win Rate	Length
GPT-4 Turbo	97.70%	2049
XwinLM 70b V0.1	95.57%	1775
PairRM+Tulu 2+DPO 70B (best-of-16)	95.40%	1607
GPT-4	95.28%	1365
Tulu 2+DPO 70B	95.03%	1418
Yi 34B Chat	94.08%	2123
PairRM+Zephyr 7B Beta (best-of-16)	93.41%	1487
LLaMA2 Chat 70B	92.66%	1790
UltraLM 13B V2.0 (best-of-16)	92.30%	1720
XwinLM 13b V0.1	91.76%	1894
UltraLM 13B (best-of-16)	91.54%	1980
Claude 2	91.36%	1069
PairRM+Tulu 2+DPO 13B (best-of-16)	91.06%	1454
Cohere Command	90.62%	1983
Zephyr 7B Beta	90.60%	1444

Aside: AlpacaEval 2

- Compare to GPT4 rather than Davinci003 (InstructGPT variant)
- Potentially too challenging to trust results
- Linear length correlation penalty is decent correction, but not a long term solution

arXiv:2404.04475v1 [cs.LG] 6 Apr 2024

Length-Controlled AlpacaEval: A Simple Way to Debias Automatic Evaluators


Yann Dubois¹, Balázs Galambosi², Percy Liang¹ and Tatsunori B. Hashimoto¹

¹Stanford University ²Independent Researcher

Abstract

LLM-based auto-annotators have become a key component of the LLM development process due to their cost-effectiveness and scalability compared to human-based evaluation. However, these auto-annotators can introduce complex biases that are hard to remove. Even simple, known confounders such as preference for longer outputs remains in existing automated evaluation metrics. We propose a simple regression analysis approach for controlling biases in auto-evaluations. As a real case study, we focus on reducing the length bias of AlpacaEval, a fast and affordable benchmark for chat LLMs that uses LLMs to estimate response quality. Despite being highly correlated with human preferences, AlpacaEval is known to favor models that generate longer outputs. We introduce a length-controlled AlpacaEval that aims to answer the counterfactual question: "What would the preference be if the model's and baseline's output had the same length?" To achieve this, we first fit a generalized linear model to predict the biased output of interest (auto-annotator preferences) based on the mediators we want to control for (length difference) and other relevant features. We then obtain length-controlled preferences by predicting preferences while conditioning the GLM with a zero difference in lengths. Length-controlling not only improves the robustness of the metric to manipulations in model verbosity, we also find that it increases the Spearman correlation with LMSYS' Chatbot Arena from 0.94 to 0.98. We release the code and resulting leaderboard.

1 Introduction



MT Bench

LLM-as-a-judge: ask a LLM (GPT4/Claude) to rate a model response:

- Two turns (response & follow-up)
- 7 categories (writing, role-play, math, coding, extraction, STEM, humanities)
- Rate one model at a time 0-10 rating scale to mitigate positional bias

MT Bench

LLM-as-a-judge: ask a LLM (GPT4/Claude) to rate a model response:

- Two turns (response & follow-up)
- 7 categories (writing, role-play, math, coding, extraction, STEM, humanities)
- 0-10 rating scale


Shortcomings: ***hard to use as sole focus during training***

- Variance in scoring up to ~0.5 points, big deltas needed for signal (via generation temperature and model API variation)
- Only 80 prompts in the eval. set
- Scoring saturated at top end (GPT4: 8.99)

Open LLM Leaderboard

Started as an engineering tool for automatically evaluating competitive models. Turned into an product with an entire team.

- Evaluate almost any model on the hub on core LLM tasks.
- Good for discovering models
- Bad for LLM developers to fixate on
- RLHF has not been shown to improve these metrics deeply, *starting to get better in 2024*




Establishing evaluation

How easy to use are these evaluations

1. ChatBotArena: **Hard to use as training signal (slow feedback), yet most reliable**
2. AlpacaEval: **Slightly expensive, for academics, training tool (~\$5 per model eval), decent correlation**
3. MT Bench: **Cheap training tool (\$.5 per model eval), decent correlation**
4. Open LLM Leaderboard: **Not super useful for studying alignment**

Chapter 3: Getting RLHF to work



Review: RLHF objective

π : LLM policy

π_θ : base LLM

x : prompt

y : completion

$$\max_{\pi_\theta} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta(y|x)} [r_\phi(x, y)] - \beta \mathbb{D}_{\text{KL}}[\pi_\theta(y | x) || \pi_{\text{ref}}(y | x)]$$

Review: RLHF objective

π : LLM policy

π_θ : base LLM

x : prompt

y : completion

$$\max_{\pi_\theta} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta(y|x)} [\underbrace{r_\phi(x, y)}] - \beta \underbrace{\mathbb{D}_{\text{KL}}[\pi_\theta(y | x) || \pi_{\text{ref}}(y | x)]}$$

Optimize “reward” *inspired* ▲
by human preferences

▲ Constrain the model to not
trust the reward too much
(preferences are hard to
model)

Review: RLHF objective

π : LLM policy

π_θ : base LLM

x : prompt

y : completion

$$\max_{\pi_\theta} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta(y|x)} [\underbrace{r_\phi(x, y)}] - \underbrace{\beta \mathbb{D}_{\text{KL}}[\pi_\theta(y | x) || \pi_{\text{ref}}(y | x)]}$$

Optimize “reward” *inspired* ▲
by human preferences

▲ Constrain the model to not
trust the reward too much
(preferences are hard to
model)

Primary questions:

1. How to implement reward: $r(x, y)$
2. How to optimize reward

Review: Preference (reward) modeling

Can we just use supervised learning on scores?

- Assigning a scalar reward of how good a response is did not work
- Pairwise preferences are easy to collect and worked!

Key idea:
Probability $\propto$ reward

Chosen completion

Prompt

Rejected completion

Score from optimal reward model

$$p^*(y_1 \succ y_2 \mid x) = \frac{\exp(r^*(x, y_1))}{\exp(r^*(x, y_1)) + \exp(r^*(x, y_2))}$$

Bradley Terry model:
Estimate probability that a given pairwise preference is true

What if we just use gradient ascent on this equation?

$$\max_{\pi_{\theta}} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_{\theta}(y|x)} [r_{\phi}(x, y)] - \beta \mathbb{D}_{\text{KL}}[\pi_{\theta}(y | x) || \pi_{\text{ref}}(y | x)]$$

What if we just use gradient ascent on this equation?

$$\max_{\pi_{\theta}} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_{\theta}(y|x)} [r_{\phi}(x, y)] - \beta \mathbb{D}_{\text{KL}}[\pi_{\theta}(y|x) \parallel \pi_{\text{ref}}(y|x)]$$

The answer, with some math, is:
Direct Preference Optimization (DPO)

Released on May 29th 2023
(4+ months before models we're discussing)





Figure 1: **DPO optimizes for human preferences while avoiding reinforcement learning.** Existing methods for fine-tuning language models with human feedback first fit a reward model to a dataset of prompts and human preferences over pairs of responses, and then use RL to find a policy that maximizes the learned reward. In contrast, DPO directly optimizes for the policy best satisfying the preferences with a simple classification objective, fitting an *implicit* reward model whose corresponding optimal policy can be extracted in closed form.



DPO core facts

1. Extremely **simple** to implement
2. **Scales nicely** with existing distributed training libraries
3. Trains an implicit reward function (can still be used as a reward model, see [RewardBench](#))

The first 2 points mean we'll see more DPO models than anything else and learn it's limits!

```
import torch.nn.functional as F

def dpo_loss(pi_logps, ref_logps, yw_idx, yl_idx, beta):
    """
    pi_logps: policy logprobs, shape (B,)
    ref_logps: reference model logprobs, shape (B,)
    yw_idx: preferred completion indices in [0, B-1], shape (T,)
    yl_idx: dispreferred completion indices in [0, B-1], shape (T,)
    beta: temperature controlling strength of KL penalty
    Each pair of (yw_idx[i], yl_idx[i]) represents the
    indices of a single preference pair.
    """

    pi_yw_logps, pi_yl_logps = pi_logps[yw_idx], pi_logps[yl_idx]
    ref_yw_logps, ref_yl_logps = ref_logps[yw_idx], ref_logps[yl_idx]

    pi_logratios = pi_yw_logps - pi_yl_logps
    ref_logratios = ref_yw_logps - ref_yl_logps

    losses = -F.logsigmoid(beta * (pi_logratios - ref_logratios))
    rewards = beta * (pi_logps - ref_logps).detach()

    return losses, rewards
```

Example code.

Rafailov, Sharma, Mitchell et al. 2023

DPO vs RL (PPO, REINFORCE, ...)

DPO and PPO are very different optimizers.

It is learning directly from preferences vs. using RL update rules.

It is also not really online vs offline RL, but that is more muddled.


More discussion:

https://twitter.com/srush_nlp/status/1729896568956895370,

<https://www.interconnects.ai/p/the-dpo-debate>,

<https://www.youtube.com/watch?v=YJMCSVLRUNs>

LEARNING FROM HUMAN FEEDBACK



Credit Tom Goldstein

<https://twitter.com/tomgoldsteins>

Aligning open language models | Lambert: 66

RLHF phase: Zephyr β

- First model to make a splash with DPO!
- Fine-tune of Mistral 7B with UltraFeedback dataset.
- Discovered weird low learning rates that are now standard ($\sim 5E-7$)
- MT Bench 7.34



RLHF phase: Tulu 2

- First model to scale DPO to 70 billion parameters!
- Strongly validated the Zephyr results.
- Started the DPO vs. PPO debate for real.
- MT Bench 70B: 7.89

Tulu v2

Open instruction & RLHF models

A12





RLHF phase: SteerLM & Starling

Still plenty of models showing that PPO (and RL methods) outperforms DPO!

- SteerLM: Attribute conditioned fine-tuning
- Starling: Introduced new preference dataset, [Nectar](#), and k-wise reward model loss function (i.e. moving beyond pairwise preferences)
 - MT Bench 7B: 8.09 (beat every model except GPT-4 at the time)

Chapter 4: Modern ecosystem




Diversity of models and more players

Examples:

- [Genstruct](#) from NousResearch: *model for rephrasing any text into instructions*
- [OLMo-Instruct](#) from AI2: truly open-source models.
- More players such as Databricks [DBRX](#) and [Cohere's Command R+](#) (first open model to pass GPT-4 on ChatBotArena)
- Research models such as Microsoft [Rho](#) (reward model weighted pretraining)
- Multilingual fine-tuning with [Aya](#)
- More MoE models ([JetMoE](#), [Qwen Moe](#),...)
- State-space models such as [Jamba](#)

Llama 3

More about scaling than alignment.
TBD if they solved the Llama 2 refusals problem.




More here: www.interconnects.ai/p/llama-3-and-scaling-open-llms

Current directions

Aligning open language models

Open vs. closed aligned models

[Image credit:](#)
Maxime Labonne



Current directions

1. **Data! Data! Data!** We are *severely limited* on experimentation by having too few preference datasets (Anthropic HH, UltraFeedback, and Nectar are main three).
2. **Continuing to improve DPO:** *tons* of papers iterating on the method ([ORPO](#), [cDPO](#), [IPO](#), [BCO](#), [KTO](#), [DNO](#), [sDPO](#), etc)
3. **More model sizes:** Most alignment research happened at 7 or 13B parameter scale. Expand up and down!
4. **Specific evaluations:** How do we get more specific evaluations than ChatBotArena?
5. **Personalization:** A large motivation behind local models, young area academically

Where open alignment is happening


- [AI2](#) (self bias): Tulu models, OLMo-Adapt, dataset releases
- [HuggingFaceH4](#): Quick releases on new base models, recipes for new techniques (e.g. ORPO / CAI), other tools
- [Berkeley-Nest/Nexusflow](#): Nectar dataset / Starling models
- [NousResearch](#): Hermes fine-tuning models, datasets, and other
- [OpenBMB](#): Preference datasets, reward models, and more
- [Argilla](#): Open preference datasets and resulting models
- Some HuggingFace users
 - [Maxime Labonne](#): Model merging & other fine-tunes
 - [Jon Durbin](#): More model merges & other fine-tunes

Thank you! Questions

Contact: nathan at natolambert dot com

Socials: @natolambert

Writing: interconnects.ai



Thanks to many teammates at HuggingFace and AI2 for supporting this journey!